

Network Layer

Dr. A Krishna Chaitanya,
Indian Institute of Information Technology Sri City

- **Functionalities**

- forwarding
- routing
- connection setup

- **Services**

- guaranteed delivery
- guaranteed delivery with bounded delay
- in-order packet delivery
- guaranteed maximum jitter

Virtual-Circuit and Datagram Networks

- **Virtual-Circuit:** provides a connection-oriented service
 - a path
 - VC numbers
 - entries in the forwarding table corresponding to each VC
- **Datagram Networks:** connectionless service
 - routers forwards packets based on destination address range or following prefix matching rule

Inside a Router





Figure 4.6 ♦ Router architecture

- Switching via memory
- Switching via bus
- Switching via interconnection of network


Scheduling and Policing

- Scheduling
 - FIFO
 - Priority Queue
 - Weighted Fair Queuing (WFQ)
- Policing
 - Leaky bucket


FIFO or FCFS




FIFO



Priority Queuing




Priority Queuing




Round Robin

- Round robin queuing discipline
- No strict priority, schedule different queues in a round robin manner.
- Work-conserving round robin discipline



Robin.jpeg

Weighted Fair Queuing



- Bandwidth - R packets per second
- Class i will get a fraction of BW equal to $\frac{w_i}{\sum_j w_j}$


Policing: Leaky Bucket

- Restrictions:
 - average rate
 - peak rate
 - burst size

Policing: Leaky Bucket


- Restrictions:
 - average rate
 - peak rate
 - burst size
- Leaky bucket:
 - a leaky bucket contains a maximum of b tokens
 - tokens are added to the bucket at rate r tokens per second
 - To transmit a packet, first remove token from the bucket and then transmit.

Leaky Bucket



- Maximum number of packets in an interval of t seconds:
 $rt + b$.

Leaky Bucket with WFQ



- Consider flow 1: Its BW is $R \frac{w_1}{\sum w_j}$

Maximum delay

- Consider flow 1: Its BW is $R \frac{w_1}{\sum w_j}$
- A burst of b_1 packets have arrived.

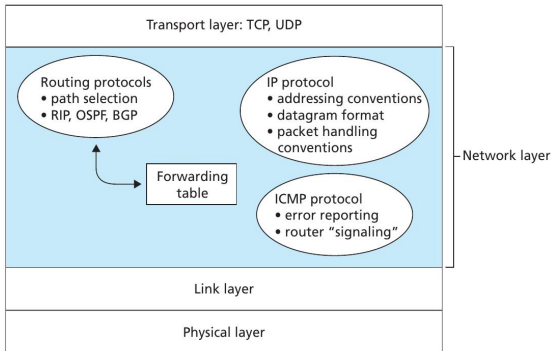
Maximum delay

- Consider flow 1: Its BW is $R \frac{w_1}{\sum w_j}$
- A burst of b_1 packets have arrived.
- Last packet will be served with in d_{max}

Maximum delay

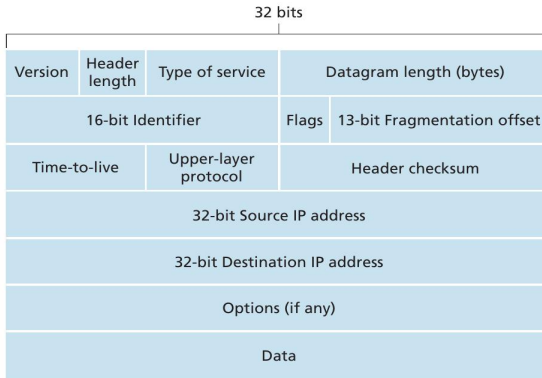
- Consider flow 1: Its BW is $R \frac{w_1}{\sum w_j}$
- A burst of b_1 packets have arrived.
- Last packet will be served with in d_{max}
- $d_{max} = \frac{b_1}{R \frac{w_1}{\sum w_j}}$

Internet Protocol



- IPv4 and IPv6

IPv4 Datagram



- Header checksum
 - needs to be computed at every router
 - TCP already has checksum, **why do datagrams need checksum?**

IPv4 Datagram

- Header checksum
 - needs to be computed at every router
 - TCP already has checksum, **why do datagrams need checksum?**
- Options
 - used for testing, debugging and security.
 - not used in general
 - Packet header is variable based on options used


IPv4 Datagram

- Header checksum
 - needs to be computed at every router
 - TCP already has checksum, **why do datagrams need checksum?**
- Options
 - used for testing, debugging and security.
 - not used in general
 - Packet header is variable based on options used
- Data: Typically TCP/UDP segment.

Fragmentation

- Datagrams are often larger than MTU
- **Fragmentation**: Splitting a larger datagram into smaller frames suitable for transmission
- Link-layer in the end system **reassembles** the fragments and forwards to network layer.
- Datagram fields:
 - **identification**: all fragments of a datagram have same identification number.
 - **flag**: indicates whether it is last fragment or not
 - **fragmentation offset**: specifies the location of the fragment in the datagram

Fragmentation



Fragmentation

Fragment	Bytes	ID	Offset	Flag
1st fragment	1,480 bytes in the data field of the IP datagram	identification = 777	offset = 0 (meaning the data should be inserted beginning at byte 0)	flag = 1 (meaning there is more)
2nd fragment	1,480 bytes of data	identification = 777	offset = 185 (meaning the data should be inserted beginning at byte 1,480. Note that $185 \cdot 8 = 1,480$)	flag = 1 (meaning there is more)
3rd fragment	1,020 bytes (= 3,980–1,480–1,480) of data	identification = 777	offset = 370 (meaning the data should be inserted beginning at byte 2,960. Note that $370 \cdot 8 = 2,960$)	flag = 0 (meaning this is the last fragment)

- Who should be IP addressed?

- Who should be IP addressed?
 - Hosts?

- Who should be IP addressed?
 - Hosts?
 - Routers?

- Who should be IP addressed?
 - Hosts?
 - Routers?
- Hosts are connected to internet through a link.
- Interface: The boundary between host and physical link.

- Who should be IP addressed?
 - Hosts?
 - Routers?
- Hosts are connected to internet through a link.
- Interface: The boundary between host and physical link.
- It is the interface that will be IP addressed.

IP Address

- IP address is **32-bit** long
- It is represented in dotted-decimal notation. Example, the address

11000001 00100000 11011000 00001001

will be represented by **193.32.216.9**

- About 4 billion addresses available
- Who assigns IP addresses?

IP Address

- IP address is 32-bit long
- It is represented in dotted-decimal notation. Example, the address

11000001 00100000 11011000 00001001

will be represented by 193.32.216.9

- About 4 billion addresses available
- Who assigns IP addresses?
- International Corporation for Assigned Names and Numbers (ICANN)

IP Address

- IP address is 32-bit long
- It is represented in dotted-decimal notation. Example, the address

11000001 00100000 11011000 00001001

will be represented by 193.32.216.9

- About 4 billion addresses available
- Who assigns IP addresses?
- International Corporation for Assigned Names and Numbers (ICANN)
- How to assign IP addresses?

IP Address


- IP address is **32-bit** long
- It is represented in dotted-decimal notation. Example, the address

11000001 00100000 11011000 00001001


will be represented by **193.32.216.9**

- About 4 billion addresses available
- Who assigns IP addresses?
- International Corporation for Assigned Names and Numbers (ICANN)
- How to assign IP addresses?
- **Subnet**: Detach each interface from its host or router, creating islands of isolated networks. Each of these isolated networks is called a subnet.

Subnet



Subnet



- The internet's addressing strategy is known as **Classless Interdomain Routing (CIDR)**
- IP broadcast address: **255.255.255.255**
- Classful addressing:
 - Class A: a.b.c.d/**8**
 - Class B: a.b.c.d/**16**
 - Class C: a.b.c.d/**24**
- CIDR: a.b.c.d/**x**

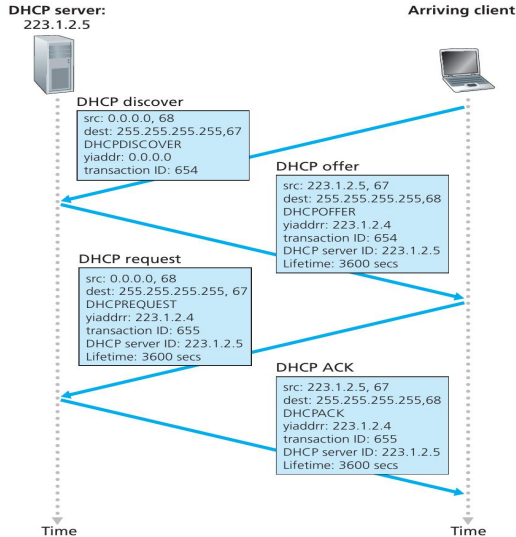
Obtaining a Block of Addresses

ISP's block	200.23.16.0/20	<u>11001000 00010111 00010000</u> 00000000
Organization 0	200.23.16.0/23	<u>11001000 00010111 00010000</u> 00000000
Organization 1	200.23.18.0/23	<u>11001000 00010111 00010010</u> 00000000
Organization 2	200.23.20.0/23	<u>11001000 00010111 00010100</u> 00000000
...
Organization 7	200.23.30.0/23	<u>11001000 00010111 00011110</u> 00000000

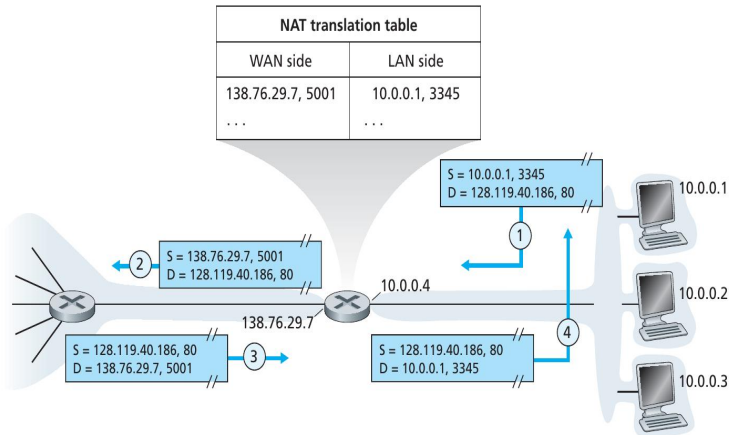
Dynamic Host Configuration Protocol

- Allows hosts to obtain IP address automatically
- Also known as plug and play protocol
- Client-server protocol
- Each server may have DHCP server. If a subnet does not have DHCP server, it will have DHCP relay agent that knows the address of DHCP server.

DHCP



Network Address Translation (NAT)



Internet Control Message Protocol (ICMP)


- Typically used for **error reporting**. Example: "Destination network unreachable"
- Can be used for congestion control
- ICMP messages have:
 - type and code field
 - header and the first 8 bytes of IP datagram that caused the ICMP message

ICMP

ICMP Type	Code	Description
0	0	echo reply (to ping)
3	0	destination network unreachable
3	1	destination host unreachable
3	2	destination protocol unreachable
3	3	destination port unreachable
3	6	destination network unknown
3	7	destination host unknown
4	0	source quench (congestion control)
8	0	echo request
9	0	router advertisement
10	0	router discovery
11	0	TTL expired
12	0	IP header bad


- Internet Engineering Task Force (IETF) developed IPv6
- Expanded addressing capabilities: 128 bits
- A streamlined 40-byte header: fixed length header
- Flow labeling and priority:
 - labeling of packets belonging to particular flows
 - ICMP packets can be given high priority than IP datagrams.
- No fragmentation and reassembly at router
- No checksum computation, and no options field.

IPv6 Datagram



- IPv4 to IPv6
 - Dual-Stack approach
 - Tunneling

Dual-Stack Approach

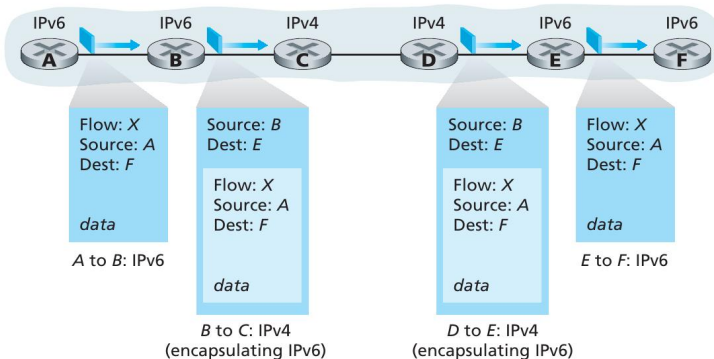


Tunneling

Logical view




Physical view



Routing

Notation



- We represent a network by an **undirected graph** $G = (N, E)$
- N is the set of Nodes (**routers**)
- E is the set of edges connecting nodes (**links**)
- $c(x, y)$ is the cost of the edge between x and y .
- Cost of a path (x_1, \dots, x_p) is sum of costs of edges along the path: $c(x_1, x_2) + \dots + c(x_{p-1}, x_p)$
- We aim to find paths with **least cost**.

- **Global** vs **Decentralized**:
 - Global routing algorithm: requires global information about links and costs at every router. Also known as Link-State algorithm
 - Decentralized routing algorithm: no node has complete information
- **Static** vs **Dynamic** routing
- **Load-sensitive** vs **Load-insensitive** routing

Link-State Routing Algorithm

- We study **Dijkstra's algorithm**
- $D(v)$: cost of the least cost path from source to destination v as of this iteration
- $p(v)$: previous node along the current least cost path from the source to v
- N' : subset of N . If $v \in N$, then least cost path to v from source is definitely known.

LS Algorithm

```
1  Initialization:
2     $N' = \{u\}$ 
3    for all nodes  $v$ 
4      if  $v$  is a neighbor of  $u$ 
5        then  $D(v) = c(u,v)$ 
6      else  $D(v) = \infty$ 
7
8  Loop
9    find  $w$  not in  $N'$  such that  $D(w)$  is a minimum
10   add  $w$  to  $N'$ 
11   update  $D(v)$  for each neighbor  $v$  of  $w$  and not in  $N'$ :
12      $D(v) = \min( D(v), D(w) + c(w,v) )$ 
13   /* new cost to  $v$  is either old cost to  $v$  or known
14     least path cost to  $w$  plus cost from  $w$  to  $v$  */
15 until  $N' = N$ 
```


LS Routing Algorithm: Example

<i>step</i>	<i>N'</i>	$D(v),p(v)$	$D(w),p(w)$	$D(x),p(x)$	$D(y),p(y)$	$D(z),p(z)$
0	u	2,u	5,u	1,u	∞	∞
1	ux	2,u	4,x		2,x	∞
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					


LS Routing: Pathology

- Congestion-sensitive routing
- Link-costs are equal to the load carried on the link.
- Link costs are not symmetric: $c(u, v) \neq c(v, u)$
- $c(u, v) = c(v, u)$, if load on the link in both directions is same


LS Routing: Pathology




a. Initial routing



b. x, y detect better path to w, clockwise



c. x, y, z detect better path to w, counterclockwise



d. x, y, z detect better path to w, clockwise

Distance-Vector (DV) Routing Algorithm

- Decentralized, asynchronous
- Iterative process
- $d_x(y)$ denotes cost of least cost path from x to y
- **Bellman-Ford** equation

$$d_x(y) = \min_v \{c(x, v) + d_v(y)\},$$

v is a neighbor of x .


- Each node x maintains the following **routing information**:
 - For each neighbor v , the cost $c(x, v)$
 - Node x 's distance vector, $\mathbf{D}_x = [D_x(y) : y \in N]$
 - Distance vectors of each of its neighbors \mathbf{D}_v

DV Algorithm

At each node, x :

```
1  Initialization:
2      for all destinations  $y$  in  $N$ :
3           $D_x(y) = c(x,y)$  /* if  $y$  is not a neighbor then  $c(x,y) = \infty$  */
4      for each neighbor  $w$ 
5           $D_w(y) = ?$  for all destinations  $y$  in  $N$ 
6      for each neighbor  $w$ 
7          send distance vector  $D_x = [D_x(y): y \text{ in } N]$  to  $w$ 
8
9  loop
10     wait (until I see a link cost change to some neighbor  $w$  or
11           until I receive a distance vector from some neighbor  $w$ )
12
13     for each  $y$  in  $N$ :
14          $D_x(y) = \min_v \{c(x,v) + D_v(y)\}$ 
15
16     if  $D_x(y)$  changed for any destination  $y$ 
17         send distance vector  $D_x = [D_x(y): y \text{ in } N]$  to all neighbors
18
19 forever
```

DV Example



Node x table

from	cost to			
	x	y	z	
x	0	2	7	
y	∞	∞	∞	
z	∞	∞	∞	

from	cost to			
	x	y	z	
x	0	2	3	
y	2	0	1	
z	7	1	0	

from	cost to			
	x	y	z	
x	0	2	3	
y	2	0	1	
z	3	1	0	

Node y table

from	cost to			
	x	y	z	
x	∞	∞	∞	
y	2	0	1	
z	∞	∞	∞	

from	cost to			
	x	y	z	
x	0	2	7	
y	2	0	1	
z	7	1	0	

from	cost to			
	x	y	z	
x	0	2	3	
y	2	0	1	
z	3	1	0	

Node z table


from	cost to			
	x	y	z	
x	∞	∞	∞	
y	∞	∞	∞	
z	7	1	0	

from	cost to			
	x	y	z	
x	0	2	7	
y	2	0	1	
z	3	1	0	

from	cost to			
	x	y	z	
x	0	2	3	
y	2	0	1	
z	3	1	0	


Time

DV Algorithm: Link Cost Changes



- Focus on **distance tables entries of y and z to x**
- At t_0 , cost has changed to 1 from 4. y updates its table with **$D_y(x) = 1$** and informs z
- At t_1 , z receives update from y and updates its table **$D_z(x) = 2$**
- At t_2 , y receives update from z and no changes in table.


DV Algorithm: Link Cost Changes



- Before link cost changes:

$$D_y(x) = 4, D_y(z) = 1, D_z(y) = 1, D_z(x) = 5$$

DV Algorithm: Link Cost Changes




- Before link cost changes:

$$D_y(x) = 4, D_y(z) = 1, D_z(y) = 1, D_z(x) = 5$$

- At t_0 , cost has changed to 60 from 4. y updates its table with

$$D_y(x) = \min\{c(y, x) + D_x(x), c(y, z) + D_z(x)\} \quad (1)$$

DV Algorithm: Link Cost Changes



- Before link cost changes:


$$D_y(x) = 4, D_y(z) = 1, D_z(y) = 1, D_z(x) = 5$$

- At t_0 , cost has changed to 60 from 4. y updates its table with

$$D_y(x) = \min\{c(y, x) + D_x(x), c(y, z) + D_z(x)\} \quad (1)$$

- $D_y(x) = \min\{60 + 0, 1 + 5\} = 6$

DV Algorithm: Link Cost Changes



- Before link cost changes:

$$D_y(x) = 4, D_y(z) = 1, D_z(y) = 1, D_z(x) = 5$$

- At t_0 , cost has changed to 60 from 4. y updates its table with


$$D_y(x) = \min\{c(y, x) + D_x(x), c(y, z) + D_z(x)\} \quad (1)$$

- $D_y(x) = \min\{60 + 0, 1 + 5\} = 6$

- At t_1 , z receives update from y and updates its table

$$D_z(x) = \min\{50 + 0, 1 + 6\} = 7$$

DV Algorithm: Link Cost Changes



- Before link cost changes:


$$D_y(x) = 4, D_y(z) = 1, D_z(y) = 1, D_z(x) = 5$$

- At t_0 , cost has changed to 60 from 4. y updates its table with

$$D_y(x) = \min\{c(y, x) + D_x(x), c(y, z) + D_z(x)\} \quad (1)$$

- $D_y(x) = \min\{60 + 0, 1 + 5\} = 6$
- At t_1 , z receives update from y and updates its table
 $D_z(x) = \min\{50 + 0, 1 + 6\} = 7$
- At t_2 , y receives update from z and updates table as
 $D_y(x) = 8$. and this process repeats.

DV Algorithm: Link Cost Changes



- Before link cost changes:

$$D_y(x) = 4, D_y(z) = 1, D_z(y) = 1, D_z(x) = 5$$

- At t_0 , cost has changed to 60 from 4. y updates its table with

$$D_y(x) = \min\{c(y, x) + D_x(x), c(y, z) + D_z(x)\} \quad (1)$$

- $D_y(x) = \min\{60 + 0, 1 + 5\} = 6$
- At t_1 , z receives update from y and updates its table
 $D_z(x) = \min\{50 + 0, 1 + 6\} = 7$
- At t_2 , y receives update from z and updates table as
 $D_y(x) = 8$. and this process repeats.
- **Count-to-infinity!**

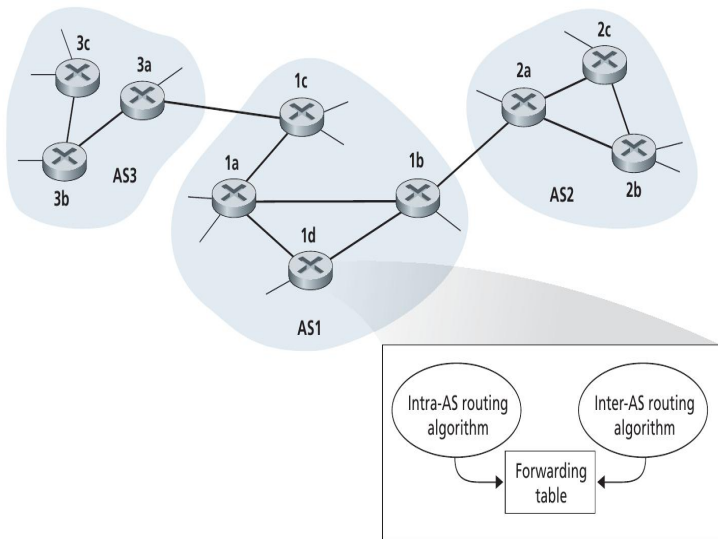
Poisoned Reverse

- If z routes through y , z will inform y that $D_z(x) = \infty$
- y cannot route to x via z as there is **no path!**
- When $c(x, y) = 60$, y updates its table with $D_y(x) = 60$!
- After receiving an update z routes to x via direct path and updates its table with $D_z(x) = 50$
- After receiving update from z , y recomputes route to x via z and informs z with $D_y(x) = \infty$ (**infact it is 51!**)

Hierarchical Routing

- **Scale**: number of routers in internet is very large. Which algorithm to use?
- **Administrative autonomy**: an organization should be able to run and administer its network as it wishes.
- These problems can be solved by organizing routers into **autonomous systems (AS)**
- Each AS will have a **gateway router**

Hierarchical routing



- Intra-AS routing
 - Routing information protocol (RIP): based on DV algorithm
 - Open shortest path first (OSPF): based on LS algorithm
- Inter-AS routing
 - Border Gateway Protocol (BGP)

Details of these algorithms (Section 4.6) are left for self study!
These are part of our CCN course