



Fundamentals of Distributed Computing

Course: Distributed Computing

Faculty: Dr. Rajendra Prasath

Spring 2019

About this Course

This course covers essential aspects that every serious programmer needs to know about **distributed algorithms, design principles and their analysis**, with emphasis on **real-time implementations** and **scalable applications development**

2

What do we learn?

- **Some Important aspects of DC:**
 - Reliable network
 - Zero Latency
 - Infinite Bandwidth
 - Secure network
 - Fixed Topology,
 - Only one administrator
 - Zero Transport cost
 - Homogeneous Network
- Remember these points when you first develop a distributed application

Reliable Network

- Hardware may fail! Power failures; Switches have a mean time between failures

The implications:

- Hardware: weight the risks of failure versus the required investment to build redundancy
- Software: we need reliable messaging: be prepared to retry messages, acknowledge messages, reorder messages (do not depend on message order), verify message integrity, and so on.

Zero Latency

Latency (not bandwidth): how much time do the data take to move from one place to another? → measures in time

- The minimum round-trip time between two points on earth is determined by the maximum speed of information transmission: the speed of light.
- At 300,000 km/sec, it will take at least 30msec to send a ping from Europe to the USA and back.

The implications:

- Strive to make as few calls over the network (other than LANs) as possible

Infinite Bandwidth

Bandwidth: how much data you can transfer over a period of time (may be measured in bits/second)

- It constantly grows
- Bandwidth may be lowered by packet loss: we may want to use larger packet sizes

The implications:

- Compression: simulate the environment to get an estimate for your needs

Secure Network

How to secure the underlying network?

The implications:

- You may need to build security into your applications from the beginning
- As a result of security considerations, you might not be able to access networked resources, different user accounts may have different privileges, and so on
- How to solve these issues efficiently?

7

Fixed Topology

Topologies do not change as long as you are in a closed environment

- In reality, servers may be added and removed often, clients (laptops, wireless ad hoc networks) are coming and going: the topology is changing constantly

The implications:

- Do not rely on specific endpoints or routes
- Abstract the physical structure of the network: the most obvious example is DNS names as opposed to IP addresses

Who is the Administrator?

Different Administrators may be associated with the network with different degrees of expertise

- Might make it difficult to locate problems
- Coordination of upgrades: will the new version of MySQL work as before with Ruby on Rails?
- Never underestimate the 'human' factor!₉

Zero Transport Cost

- Going from the application layer to the transport layer (2nd highest in the five layer TCP/IP reference model) is not free:
- Information needs to be serialized (marshalling) to get data onto the wire
- The cost (in terms of money) from setting and running the network is not zero.
- Have we leased the necessary bandwidth

Everything cost "Some money" !! 0

Homogeneous Network

Homogeneous = of the same kind; uniform

- ➔ Even a home network may connect a Linux PC and a Windows PC. A homogeneous network today is the exception, not the rule!

The implications:

- ➔ Interoperability will be needed
- ➔ Use standard technologies such as XML

11



FUNDAMENTALS OF DISTRIBUTED SYSTEMS

Can you parallelize all apps?


- What can you parallelize?
 - Identify Instructions that could execute in parallel?
- If the proportion of an application that you can parallelize is x ($0 < x < 1$),
Maximum speed up = $1/(1-x)$
- Let us assume that a task needs ts time to run on one machine. How much time does it take to run the same task on p CPUs?
- Let running time be tp of a task running on p CPUs. Now $tp = to + ts * (1 - x + x/p)$, where to is the overhead added due to parallelisation (communication, synchronization, etc)

13


Flynn's Classification

- Single Instruction Single Data (SISD) Stream
 - Traditional von Neumann architecture
- Single Instruction Multiple Data (SIMD) Stream
 - Scientific Applications, Vector Processors, array processors and so on
- Multiple Instruction Single Data (MISD) Stream
 - Visualization is an example
- Multiple Instruction Multiple Data (MIMD) Stream
 - Distributed Systems


Flynn's Classification



(a) *SIMD*



(b) *MIMD*



(c) *MISD*


C Control Unit

P Processing Unit


I instruction stream

D data stream


Interconnection Topologies




a) Bus




b) Linear array




c) Star




d) Ring




e) Tree



f) Near-neighbor mesh




g) Completely connected



h) 3-cube (hypercube)


Interconnection Topologies (contd)

→ Wrap - around
mesh (torus)



k-ary d-cubes

→ Hypercube of Dimension 4



→ k-ary d-cubes (generalized version)

Message Passing Systems

Basic Primitive Operations

→ Send

→ *send* message from process A to Process B
 $A \rightarrow B$

→ Receive

→ *receive* message at Process B from Process A
 $B \leftarrow A$

→ Compute at A and / or B

→ do the specific computations at A and / or B

Message Passing vs. Shared Memory

→ Emulate MP on SM

- Partition the shared address space
- Send / Receive messages via shared address space

→ Emulate SM on MP

- Model each shared location as a separate process
- Send: Write to the shared object by sending messages to owner process for the object
- Receive: Read from shared object by sending query to the owner process of the object

Synchronous vs. Asynchronous

→ Synchronous (send / receive)

- Handshake between sender and receiver
- *send* completes only when *receive* completes
- *receive* completes only when copying of the data to the buffer is over

→ Asynchronous (send)

- No need for handshake between sender and receiver
- Control returns to the invoking process when data copied out of user-specified buffer

21

Blocking vs. non-blocking

→ Blocking

- Control returns to the invoking process after the the task completes

→ Non-blocking

- Control returns to the invoking process when data copied out of user-specified buffer
- *send* completes even before copying the data to the user buffer
- *receive* may happen even before data may have arrived from the sender
- How to order EVENTS?

Applications

- Mobile Systems
- Sensor networks
- Pervasive Computing
 - Smart workplace
 - Intelligent Home
- Peer-to-peer computing
- Distributed Agents
- Distributed Data Mining
- Grid Computing
- Security aspects in Distributed Systems

23

Summary

- **Focused aspects** in this course:
 - Fundamental aspects while building distributed applications
 - Interconnection topologies
 - Number of Message exchanges
 - Primitives of Distributed Communications
 - Message Passing vs Shared memory systems
 - Properties of a distributed system
 - Many more to come up ... stay tuned in ...

How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <http://www.iiits.ac.in/FacPages/index-rajendra.html>

OR

→ <http://rajendra.2power3.com>

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

Thanks ...

