

# Distributed Mutual Exclusion Algorithms

Course: Distributed Computing

Faculty: Dr. Rajendra Prasath

# About this topic

This course covers various concepts in **Mutual Exclusion in Distributed Systems**. We will also focus on different types of distributed mutual exclusion algorithms in distributed contexts and their analysis

# What did you learn so far?

- Challenges in Message Passing systems
- Distributed Sorting
- Space-Time Diagram
- Partial Ordering / Causal Ordering
- Concurrent Events
- Local Clocks and Vector Clocks
- Distributed Snapshots
- Termination Detection
- Topology Abstraction and Overlays
- Leader Election Problem in Rings
- Message Ordering / Group Communications

# Recent Topic ...

- Communication Models
- Design Issues
  - Process Failures
- Message Ordering
  - Good / Bad ordering
  - Various Types of Ordering of messages
- Group Communication
  - Causal ordering based approach
  - Many more to come up ... stay tuned in !!

# Message Ordering (recap)

## ➡ How to order messages?

- ➡ Send vs Delivery

- ➡ Global Time Ordering

- ➡ Total Ordering

- ➡ Causal Ordering

- ➡ Sync Ordering

- ➡ FIFO Ordering

- ➡ Unordered multicast

## ➡ Good / Bad Ordering

# Topics to focus on ...

- Leader Election in Distributed Systems
- Topology Abstraction and Overlays
- Message Ordering
- Group Communication

## → Distributed Mutual Exclusion

---

- Deadlock Detection
- Check pointing and rollback recovery

For Mid Semester 2

# Mutual Exclusion in Distributed Systems

Let us explore mutex algorithms proposed for various interconnection networks

# Why do we need Mutex?

## → Mutual Exclusion

### → Operating systems: Semaphores

→ In a single machine, you could use semaphores to implement mutual exclusion

→ How to implement semaphores?

→ Inhibit interrupts

→ Use clever instructions (e.g. test-and-set)

→ On a multiprocessor shared memory machine, only the latter works



# Characteristics

- Processes communicate only through messages
  - no shared memory or no global clocks
- Processes must expect unpredictable message delays
- Processes coordinate access to shared resources (printer, file, etc.) that should only be used in a mutually exclusive manner.

# Race Conditions

- Consider Online systems - For example, Airline reservation systems maintain records of available seats
- Suppose two people buy the same seat, because each checks and finds the seat available, then each buys the seat
- Overlapped accesses generate different results than serial accesses - race condition

# Distributed Mutual Exclusion

## → Needs

- Only one process should be in critical section at any point of time
- What about resources?

# Distributed Mutual Exclusion

- **No Deadlocks** - no set of sites should be permanently blocked, waiting for messages from other sites in that set
- **No starvation** - no site should have to wait indefinitely to enter its critical section, while other sites are executing the CS more than once
- **Fairness** - requests honored in the order they are made. This means processes have to be able to agree on the order of events. (Fairness prevents starvation.)
- **Fault Tolerance** - the algorithm is able to survive a failure at one or more sites

# Distributed MutEx – An overview

**Token-based solution:** Processes share a special message known as a token

- ➔ Token holder has right to access shared resource
- ➔ Wait for/ask for (depending on algorithm) token; enter Critical Section (CS) when it is obtained, pass to another process on exit or hold until requested (depending on algorithm)
- ➔ If a process receives the token and doesn't need it, just pass it on

# Distributed MutEx – A Few Issues

- Who can access the resource?
- When does a process to be privileged to access the resource?
- How long does a process access the resource?  
Any finite duration?
- How long can a process wait to be privileged?
- Computation complexity of the solution

# Types of Distributed MutEx

- Token-based distributed mutual exclusion algorithms
  - Suzuki - Kasami's Algorithm
- Non-token based distributed mutual exclusion algorithms
  - Lamport's Algorithm
  - Ricart-Agartala's Algorithm

# Token Based Methods

## Advantages:

- Starvation can be avoided by efficient organization of the processes
- Deadlock is also avoidable

## Disadvantage: Token Loss

- Must initiate a cooperative procedure to recreate the token
- Must ensure that only one token is created!



# Non-Token Based Methods

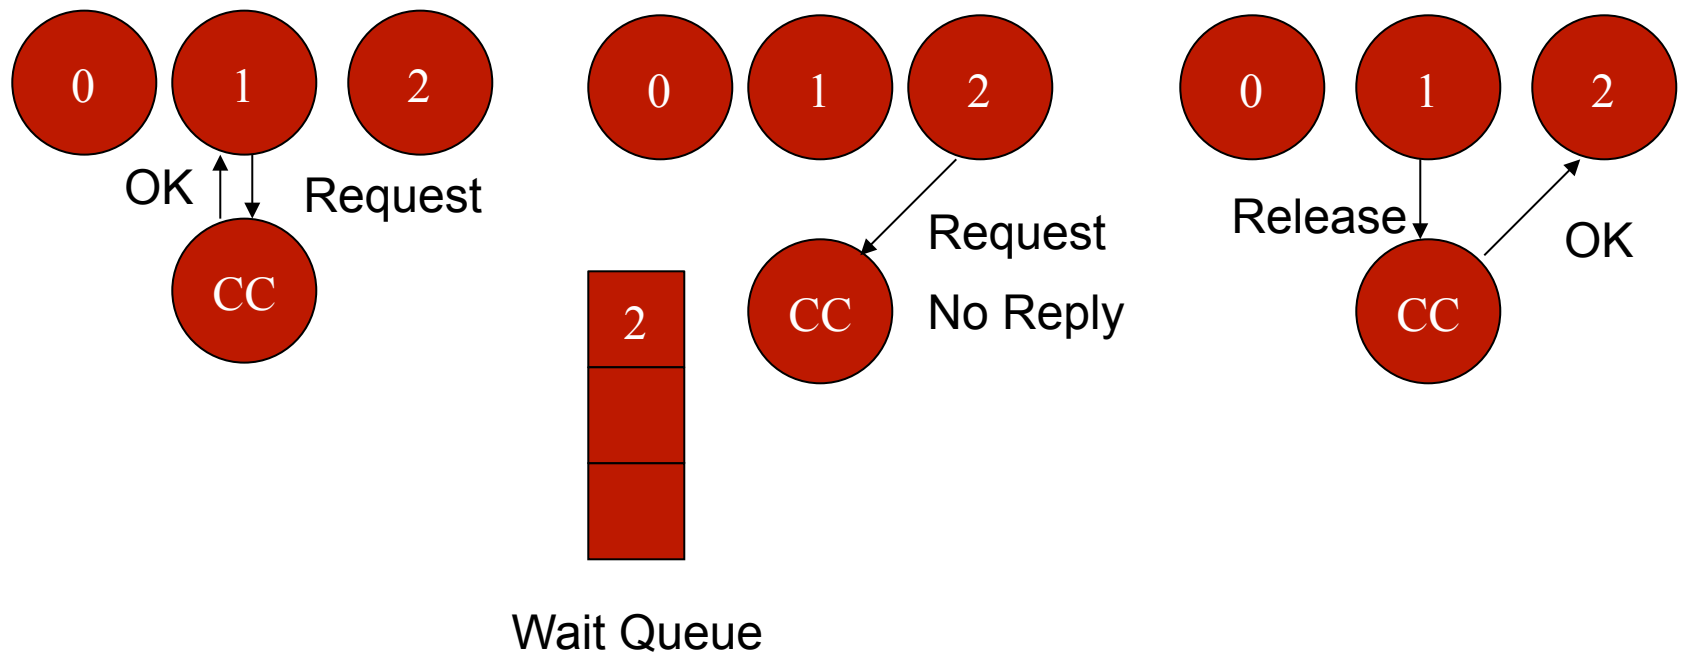
- ➔ **Permission-based solutions:** a process that wishes to access a shared resource must first get permission from one or more other processes.
- ➔ Avoids the problems of token-based solutions, but is more complicated to implement

# Basic Algorithms

- Centralized
- Decentralized
- Distributed
  - Distributed with “voting” – for increased fault tolerance
- Token Ring

# Centralized Mutual Exclusion

- Central coordinator manages the FIFO queue of requests to guarantee "no starvation"



# Performance Analysis

- Guarantees mutual exclusion
- No starvation: Only if requests served in order
- No deadlock
- Fault tolerant?
  - Single point of failure
  - Blocking requests mean client processes have difficulty distinguishing crashed coordinator from long wait
  - Bottlenecks
- The solution is simple and ease

# Centralized Control of MutEx

- A central coordinator (master or leader)
- Is elected (which algorithm?)
- Grants permission to enter CS & keeps a queue of requests to enter the CS.
- Ensures only one process at a time can access the CS
- Has a special token message, which it can give to any process to access CS

# Centralized Control - Operations

- To enter a CS, send a request to the coordinator & wait for token.
- On exiting the CS, send a message to the coordinator to release the token.
- Upon receipt of a request, if no other process has the token, the coordinator replies with the token; otherwise, the coordinator queues the request
- Upon receipt of a release message, the coordinator removes the oldest entry in the queue (if any) and replies with a token

# Centralized Control - Features

- Safety, Liveness are guaranteed
- Ordering also guaranteed (what kind?)
- Requires 2 messages for entry + 1 messages for exit operation.
- Client delay: one round trip time (request + grant)
- Synchronization delay: 2 message latencies (release + grant)
- The coordinator becomes performance bottleneck and single point of failure

# Decentralized MutEx

- More fault-tolerant than centralized approach
- Uses the Distributed Hash Table (DHT) approach to locate objects/replicas
  - Object names are hashed to find the node where they are stored (succ function)
- $n$  replicas of each object are placed on  $n$  successive nodes
  - Hash object name to get addresses
- Now every replica has a coordinator that controls access



# The Decentralized Algorithm

→ Coordinators respond to requests at once:

Yes OR No

→ **Majority:** To use the resource, a process must receive permission from  $m > n/2$  coordinators

→ If the requester gets fewer than  $m$  votes, it will wait for a random time and then ask again

→ If a request is denied, or when the CS is completed, notify the coordinators who have sent OK messages, so they can respond again to another request. (Why is this important?)

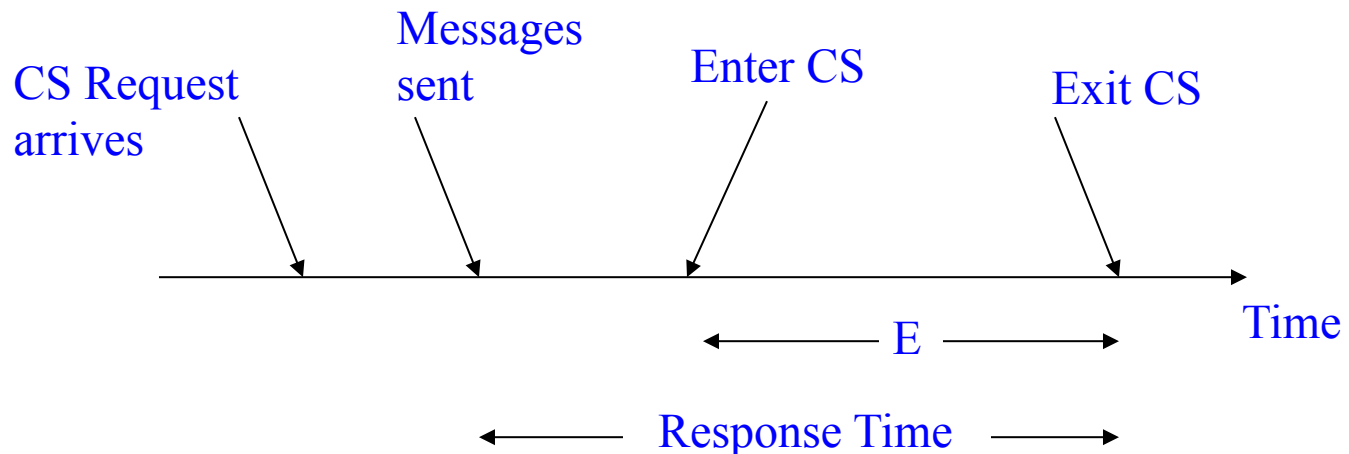

# The Decentralized Algo - Analysis

- More robust than the central coordinator approach. If one coordinator goes down others are available.
- If a coordinator fails and resets then it will not remember having granted access to one requestor, and may then give access to another.
- It is highly unlikely that this will lead to a violation of mutual exclusion

# The Decentralized Algorithm - Issues

- If a resource is in high demand, multiple requests will be generated by different processes.
- High level of contention
- Processes may wait a long time to get permission - Possibility of starvation exists
- Resource usage drops.

# Performance Metrics



# Performance - Analysis

- Number of messages per CS invocation: should be minimized
- Synchronization delay, i.e., time between the leaving of CS by a site and the entry of CS by the next one: should be minimized
- Response time: time interval between request messages transmissions and the exit of CS
- System throughput, i.e., rate at which system executes the requests for CS: should be maximized
- If **d** is the synchronization delay, **e** the average CS execution time:

$$\text{System Throughput} = 1 / (d + e)$$

# Performance - Analysis (contd)

- **Low and High Load:**
  - Low load: No more than one request at a given point in time
  - High load: Always a pending mutual exclusion request at a site
- **Best and Worst Case:**
  - Best Case (low loads): Round-trip message delay + Execution time -  $2T + E$
  - Worst case (high loads)
- Message traffic: low at low loads, high at high loads
- Average performance: when load conditions fluctuate widely

# Token Ring Approach


- Processes are organized in a logical ring:  $P_i$  has a communication channel to  $P_{(i+1) \bmod N}$

## Operations:

- Only the process holding the token T can enter the CS
- To enter the critical section, wait passively for T  
When in CS, hold on to T and don't release it
- To exit the CS, send T onto your neighbor
- If a process does not want to enter the CS when it receives T, it simply forwards T to the next neighbor

# Token Rings - Illustration

➔ Request movements in an unidirectional ring network





# Token Rings - Features

- Safety & Liveness are guaranteed
- Ordering is not guaranteed
- Bandwidth: 1 message per exit
- Client delay: 0 to N message transmissions
- Synchronization delay between one process's exit from the CS and the next process's entry is between 1 and N-1 message transmissions

# Non-Token Based Algorithms

## → Notations:

→  $P_i$ :  $i^{\text{th}}$  Process

→  $R_i$ : Request set, containing IDs of all  $P_i$  s from which permission must be received before accessing CS

→ Non-token based approaches use time stamps to order requests for CS

→ Smaller time stamps get priority over larger ones

## → Lamport's Algorithm

→  $R_i = \{P_1, P_2, \dots, P_n\}$ , i.e., all processes.

→ Request queue: maintained at each  $P_i$  ordered by time stamps.

→ **Assumption:** message delivered in FIFO

# Lamport's Algorithm

## → Requesting CS:

- Send REQUEST( $ts_i, i$ ) where ( $ts_i, i$ ) - Request time stamp; Place REQUEST in  $request\_queue_i$
- On receiving the message;  $P_j$  sends time-stamped REPLY message to  $P_i$ ;  $P_i$ 's request placed in  $request\_queue_j$

## → Executing CS:

- $P_i$  has received a message with time stamp larger than ( $ts_i, i$ ) from all other sites
- $P_i$ 's request is the top most one in  $request\_queue_i$

## → Releasing CS:

- Exiting CS: send a time stamped RELEASE message to all sites in its request set
- Receiving RELEASE message:  $P_j$  removes  $P_i$ 's request from its queue

# Notable Points

- Purpose of REPLY messages from  $i$  to  $j$  is to ensure that  $j$  knows of all requests of  $i$  prior to sending the REPLY (possibly any request of  $i$  with timestamp lower than  $j$ 's request)
- Requires FIFO channels
- $3(n - 1)$  messages per critical section invocation
- Synchronization delay = max msg transmission time
- Requests are granted in order of increasing timestamps

# Performance Improvements

- 3(n-1) messages per Critical Section invocation
  - (n - 1) REQUEST messages
  - (n - 1) REPLY messages
  - (n - 1) RELEASE messages
- Synchronization delay: T
- Optimization:
  - Suppress reply messages: For example,  $P_j$  receives a REQUEST message from  $P_i$  after sending its own REQUEST message with time stamp higher than that of  $P_i$ 's then Do NOT send a REPLY message
  - Messages reduced to between 2(n-1) and 3(n-1)

# Ricart & Agrawala's Algorithm

- A time-stamp based approach
- Originally proposed by Lamport using logical clocks
- Modified by Ricart & Agrawala

# Ricart & Agrawala's Algorithm

## Main Idea:

- Process  $j$  need not send a REPLY to Process  $i$  if  $j$  has a request with timestamp lower than the request of  $i$  (since  $i$  cannot enter before  $j$  here)
- Does not require FIFO
- $2(n - 1)$  messages per critical section invocation
- Synchronization delay = maximum message transmission time
- Requests granted in order of increasing timestamps

# Ricart & Agrawala (contd)

- Processes need entry to critical section multicast a request, and can enter it only when all other processes have replied positively
- Messages requesting entry are of the form  $\langle T, P_i \rangle$ 
  - $T$  - sender's timestamp (Lamport clock)
  - $P_i$  the sender's identity



# Ricart & Agrawala - Algorithm

To enter the Critical Section (CS):

- Set state = wanted
- multicast "request" to all processes (including timestamp)
- wait until all processes send back "reply"
- change state to held and enter the CS

On receipt of a request  $\langle T_j, P_j \rangle$  at  $P_i$ :

- if (state == held) or (state == wanted &  $(T_i, P_i) < (T_j, P_j)$ ) then enqueue the request
- else "reply" to  $P_j$

On exiting the CS:

- change state to release and "reply" to all queued requests

# Ricart & Agrawala - Simplified

To request Critical Section:

- send timestamped REQUEST message  $(ts_i, i)$

On receiving request  $(ts_i, i)$  at  $j$ :

- if  $j$  is neither requesting nor executing critical section then send REPLY to  $i$
- if  $j$  is requesting and  $i$ 's request timestamp is smaller than  $j$ 's request timestamp then
  - enqueue the request; Otherwise, defer the request

To enter Critical Section:

- Process  $i$  enters critical section on receiving REPLY messages from all processes

To release Critical Section:

- send REPLY to all deferred requests

# Summary

- Mutual Exclusion Problem
- Basics of Mutex algorithms
- Various Types of Mutex algorithms
  - Token-based
    - Token rings
  - Non-Token based algorithm
    - Lamport's Algorithm
    - Ricart - Agrawala's Algorithm
- Performance Metrics
  - Many more to come up ... stay tuned in !!

# How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <http://www.iiits.ac.in/FacPages/index-rajendra.html>


OR

→ <http://rajendra.2power3.com>

# Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
  - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

# Thanks ...



## ... Questions ???