

Check Pointing and Rollback Recovery

Course: Distributed Computing

Faculty: Dr. Rajendra Prasath

Spring 2019

About this topic

This course covers various concepts in **Check Pointing and Rollback Recovery**. We will also focus on the essential aspects of check pointing and roll back recovery in distributed contexts

What did you learn so far?

- Challenges in Message Passing systems
- Distributed Sorting
- Space-Time Diagram
- Partial Ordering / Causal Ordering
- Concurrent Events
- Local Clocks and Vector Clocks
- Distributed Snapshots
- Termination Detection
- Topology Abstraction and Overlays
- Leader Election Problem in Rings
- Message Ordering / Group Communications
- Distributed Mutual Exclusion Algorithms

Topics to focus on ...

- Distributed Mutual Exclusion
- Deadlock Detection
- Check Pointing and Rollback Recovery
- Self-Stabilization
- Distributed Consensus
- Reasoning with Knowledge
- Peer - to - peer computing and Overlays
- Authentication in Distributed Systems

For End Semester

Distributed Mutual Exclusion(Recap)

- **No Deadlocks** - No processes should be permanently blocked, waiting for messages (Resources) from other sites
- **No starvation** - no site should have to wait indefinitely to enter its critical section, while other sites are executing the CS more than once
- **Fairness** - requests honored in the order they are made. This means processes have to be able to agree on the order of events. (Fairness prevents starvation)
- **Fault Tolerance** - the algorithm is able to survive a failure at one or more sites

Deadlock - Illustrated (Recap)


→ Vehicular Traffic - A real-time scenario



Dining Philosophers (Recap)

- Each philosopher must alternately think and eat
- A philosopher can only eat when they have both left and right forks
- **Problem:** How to design a discipline of behavior (a concurrent algorithm) such that no philosopher will starve?

→ Suggest a Simple Solution ??



Check Pointing and Rollback Recovery


Let us explore Check Pointing and Roll Back Recovery algorithms in distributed systems

Handling Failures / Recovery?

- Failure of a site/node in a distributed system causes **inconsistencies** in the state of the system.
- Recovery: bringing back the failed node in step with other nodes in the system.
- Failures:
 - **Process failure:**
 - Deadlocks, protection violation, erroneous user input, etc.
 - **System failure:**
 - Failure of processor/system. System failure can have full/partial amnesia.
 - It can be a pause failure (system restarts at the same state it was in before the crash) or a complete halt.
 - **Secondary storage failure:** data inaccessible.
 - **Communication failure:** network inaccessible.


Recovery in Concurrent Systems

- State involves message exchanges in DS
- In distributed systems, rolling back one process can cause the roll back of other processes
- **Orphan messages & Domino effect:** Assume Y fails after sending m
 - X has record of m at x3 but Y has no record. M → orphan message.
 - Y rolls back to y2 → X should go to x2
 - If Z rolls back, X and Y has to go to x1 and y1 → Domino effect, roll back of one process causes one or more processes to roll back




Messages Lost

- If Y fails after receiving m, it will rollback to y1
- X will rollback to x1
- m will be a lost message as X has recorded it as sent & Y has no record of receiving it




Livelocks



- ➔ Y crashes before receiving `n1`. Y rolls back to `y1` → X to `x1`
- ➔ Y recovers, receives `n1` and sends `m2`
- ➔ X recovers, sends `n2` but has no record of sending `n1`
- ➔ Hence, Y is forced to rollback second time. X also rolls back as it has received `m2` but Y has no record of `m2`
- ➔ Above sequence can repeat indefinitely, causing a livelock

Consistent Checkpoints



- ➔ Overcoming domino effect and livelocks: checkpoints should not have messages in transit.
- ➔ Consistent checkpoints: no message exchange between any pair of processes in the set as well as outside the set during the interval spanned by checkpoints.
- ➔ $\{x1, y1, z1\}$ is a strongly consistent checkpoint

Types of CRR Algorithms

- **Synchronous Algorithm**
 - Two Phase algorithm proposed by Koo and Toueg
- **Asynchronous Algorithm**
 - A simple algorithm proposed by Juang & Venkatesan

Consistent Set of Checkpoints

Assumptions:

- Checkpoint, **send** / **recv** are atomic
- Take a checkpoint after sending every message
- The set of the most recent checkpoints is always consistent
 - Why? Is it strongly consistent?
- What is the main problem with this approach?
- Take a checkpoint after every K messages sent?
- Is it still consistent?

Synchronous Checkpointing Algo

→ Proposed by Koo and Toueg¹ (1987)

→ **Assumptions:**

- processes communicate by exchanging messages through channels
- channels are FIFO, end-to-end protocols cope up with the message loss due to rollback recovery
- Communication failures do not partition the network
- Uses two kinds of checkpoints
 - Tentative
 - Permanent

¹ R. Koo and S. Toueg, "Checkpointing and Rollback-Recovery for Distributed Systems," in IEEE Transactions on Software Engineering, vol. SE-13, no. 1, pp. 23-31, Jan. 1987. doi: 10.1109/TSE.1987.232562

Phase - 1

- Initiator: take tentative checkpoint
- Initiator requests all other processes to take tentative checkpoint
- All other processes:
 - can respond 'yes' or 'no'
- Initiator: decide to make checkpoints permanent if everyone has responded 'yes'
- A process can fail to take a checkpoint due to the nature of application (e.g.,) lack of log space, unrecoverable transactions

Phase - 2

- If all processes took checkpoints, P_i decides to make the checkpoint permanent.
- Otherwise, checkpoints are to be discarded.
- P_i conveys this decision to all the processes as to whether checkpoints are to be made permanent or to be discarded

Potential Issues

- Between tentative checkpoint and commit/abort of checkpoint process must hold back messages.
- Does this guarantee we have a strongly consistent state?
- Can you construct an example that shows we can still have lost messages?

Synchronous Checkpointing: Properties

- All or none of the processes take permanent checkpoints
- There is no record of a message being received but not sent
- Checkpoints may be taken unnecessarily (Give an example!!)
- Can these unnecessarily checkpoints be avoided?

Optimizing Checkpoints

Main IDEA:

- Record all messages sent and received after the last checkpoint ($\text{last_recv}(x, y)$, $\text{first_sent}(x, y)$)
- When X requests Y to take a tentative checkpoint:
 - X sends the last message received from Y with the request
 - Y takes a tentative checkpoint only if the last message received by X from Y was sent after Y sent the first message after the last checkpoint (**Happened before !!**)
$$\text{last_recv}(x, y) \geq \text{first_sent}(y, x)$$
- When a process takes a checkpoint, it will ask all other processes (that sent messages to the process) to take checkpoints.

Rollback Recovery: Properties

- There are two phases: Phase 1 and Phase 2
- Assume that between requests to rollback and decision, no one sends other messages
- All or none of the processes restart from checkpoints
- After rollback, all processes resume in a consistent state
- Can have unnecessary rollback: can use a similar technique as the one in taking checkpoints to eliminate unnecessary rollback

Rollback Recovery

→ Phase 1


- Initiator: check whether all processes are willing to restart from last checkpoints
- Others: may reply 'yes' or 'no'

→ Phase 2

- Initiator: propagate go/nogo decision to all processes
- Others: carry out the decision of the initiator

Unnecessary Rollbacks

- Avoid Rollback in unnecessary situations?
- An example
 - (z_2 does not need to rollback - why?)



Disadvantages


- Check Pointing Algorithm generates **message traffic**
- Synchronization delays are introduced
- These costs may seem high if failures between checkpoints are unlikely

Asynchronous Approach

- Take multiple local checkpoints independently
- After a failure, try to find a consistent set of recent checkpoints
- All incoming messages between local checkpoints are logged
 - pessimistic approach: log each message before processing
 - optimistic approach: buffer messages & log in batches
- Why is the second approach called optimistic?
- What are the advantages and disadvantages of each approach?

An Event Driven Computation

- A process waits until it receives a message; then processes the received message; changes its state and sends zero or more messages to its neighbors and then waits to receive the next message
- The current state and the contents of the messages sent depend on its previous state and the content of the message
- Events are identified by unique numbers (increasing)



Asynchronous Checkpointing Algo

→ Proposed by Juang & Venkatesan²

Assumptions:

- Communication channels are reliable
- Communication channels are FIFO
- Communication channels have no buffer size limits
- Message transmission delay is bounded
- Underlying system is **Event-Driven**, with locally timestamped (monotonically increasing numbers) events: Each event waits for a message, processes the message, changes process state, and sends a number of messages

² <https://www.utdallas.edu/~venky/pubs/crash-rec-icdcs91.pdf>

Basic Idea

- At each event, a triplet $\{s, m, \text{msgs_sent}\}$ is put in the the log: s is the state, m is the message causing the event, msgs_sent is the set of messages sent.

Two data structures used:

- $\text{RCVD}(i, j, \text{checkpoint})$ -- the number of message received by processor i from processor j at checkpoint,
- $\text{SENT}(i, j, \text{checkpoint})$ -- the number of messages sent from i to j at checkpoint.
- Use the message send/rcv counts to determine the point to rollback.

Algorithm

At process i:

- If i is a process that is recovering from a failure, checkpoint = the latest event logged in the stable storage.
- else checkpoint = latest event that took place.
- for k = 1 to N do
 - send ROLLBACK(i, SENT(i, j, checkpoint)) to all neighbors j
 - wait for ROLLBACK messages from all neighbors
 - for every ROLLBACK(j, c) received
 - if (RCVD(i, j, checkpoint) > c) then
 - find the latest event e such that RCVD(i, j, e) = c
 - checkpoint = e

Is the algorithm consistent?

→ In each iteration:

At least one processor **will rollback** to its final recovery point **unless current recovery point is consistent**

→ Answer: YES / NO

→ Complexity of this algorithm?

→ will it be greater than $O(n)$ where n is the total number of message exchanges?

→ Explore the details ... !!

Summary

- Recovery in Distributed / Concurrent Systems
- Checkpointing
 - Consistent set of checkpoints
- Rollback recovery
 - Synchronous Algorithm (Koo and Toueg)
 - Asynchronous Algorithm (Juang & Venkatesan)
- Stay tuned ... More to come up ... !!

How to reach me?

→ Please leave me an email:

rajendra [DOT] prasath [AT] iiits [DOT] in

→ Visit my homepage @

→ <http://www.iiits.ac.in/FacPages/index-rajendra.html>

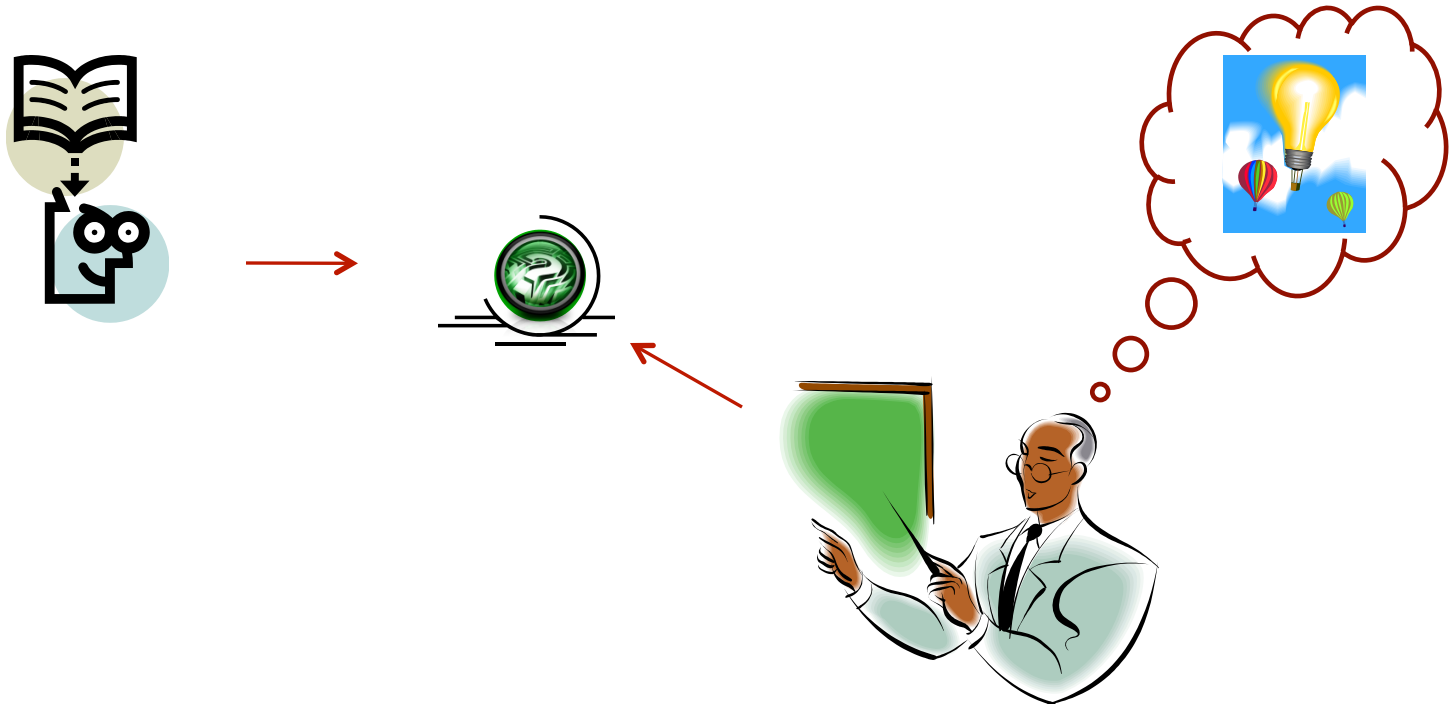
OR

→ <http://rajendra.2power3.com>

Help among Yourselves?

- **Perspective Students** (having CGPA above 8.5 and above)
- **Promising Students** (having CGPA above 6.5 and less than 8.5)
- **Needy Students** (having CGPA less than 6.5)
 - Can the above group help these students? (Your work will also be rewarded)
- You may grow a culture of **collaborative learning** by helping the needy students

Thanks ...



... Questions ???